

UL-Inclass-day-2

March 11, 2020

```
[85]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns; sns.set(style="ticks", color_codes=True)
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from scipy.stats import zscore
import sklearn.metrics
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
import warnings
warnings.filterwarnings("ignore")
```

```
[86]: df=pd.read_csv('2Classdata.csv')
df.head()
```

```
[86]:
```

	pelvic_incidence	pelvic_tilt	numeric	lumbar_lordosis_angle	sacral_slope	\
0	63.027818		22.552586	39.609117	40.475232	
1	39.056951		10.060991	25.015378	28.995960	
2	68.832021		22.218482	50.092194	46.613539	
3	69.297008		24.652878	44.311238	44.644130	
4	49.712859		9.652075	28.317406	40.060784	

	pelvic_radius	degree_spondylolisthesis	class
0	98.672917	-0.254400	Abnormal
1	114.405425	4.564259	Abnormal
2	105.985135	-3.530317	Abnormal
3	101.868495	11.211523	Abnormal
4	108.168725	7.918501	Abnormal

1 reading of data

```
[87]: df1=df.copy()
      df1.head()
```

```
[87]:   pelvic_incidence  pelvic_tilt numeric  lumbar_lordosis_angle  sacral_slope \
0         63.027818         22.552586         39.609117         40.475232
1         39.056951         10.060991         25.015378         28.995960
2         68.832021         22.218482         50.092194         46.613539
3         69.297008         24.652878         44.311238         44.644130
4         49.712859          9.652075         28.317406         40.060784

      pelvic_radius  degree_spondylolisthesis  class
0         98.672917          -0.254400  Abnormal
1        114.405425           4.564259  Abnormal
2        105.985135          -3.530317  Abnormal
3        101.868495          11.211523  Abnormal
4        108.168725           7.918501  Abnormal
```

```
[88]: df1['class']=pd.get_dummies(df['class'])
```

```
[89]: df1.head()
```

```
[89]:   pelvic_incidence  pelvic_tilt numeric  lumbar_lordosis_angle  sacral_slope \
0         63.027818         22.552586         39.609117         40.475232
1         39.056951         10.060991         25.015378         28.995960
2         68.832021         22.218482         50.092194         46.613539
3         69.297008         24.652878         44.311238         44.644130
4         49.712859          9.652075         28.317406         40.060784

      pelvic_radius  degree_spondylolisthesis  class
0         98.672917          -0.254400         1
1        114.405425           4.564259         1
2        105.985135          -3.530317         1
3        101.868495          11.211523         1
4        108.168725           7.918501         1
```

```
[90]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 310 entries, 0 to 309
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   pelvic_incidence                      310 non-null    float64
1   pelvic_tilt numeric                    310 non-null    float64
2   lumbar_lordosis_angle                  310 non-null    float64
```

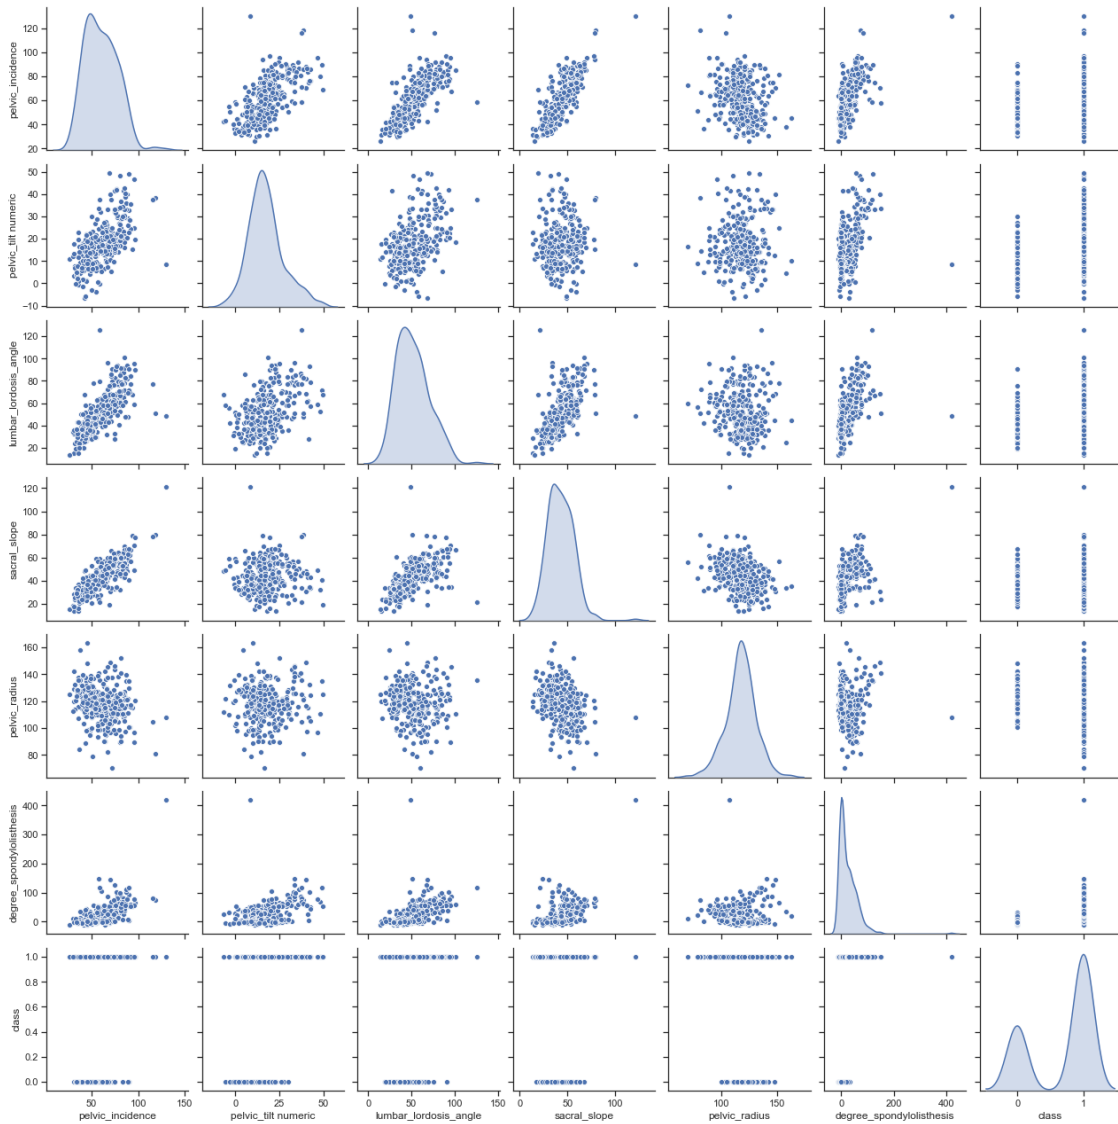
```
3  sacral_slope          310 non-null    float64
4  pelvic_radius         310 non-null    float64
5  degree_spondylolisthesis 310 non-null    float64
6  class                 310 non-null    uint8
dtypes: float64(6), uint8(1)
memory usage: 15.0 KB
```

```
[91]: df1.isnull().sum()
```

```
[91]: pelvic_incidence          0
      pelvic_tilt numeric       0
      lumbar_lordosis_angle     0
      sacral_slope              0
      pelvic_radius             0
      degree_spondylolisthesis  0
      class                    0
      dtype: int64
```


```
[92]: import seaborn as sns
      sns.pairplot(df1,diag_kind='kde')
```

```
[92]: <seaborn.axisgrid.PairGrid at 0x23bda190400>
```




```
[93]: sns.pairplot(df1,diag_kind='kde', hue='class')
```

```
[93]: <seaborn.axisgrid.PairGrid at 0x23bda184780>
```



```
[94]: sns.heatmap(df1.corr(),annot=True)
```

```
[94]: <matplotlib.axes._subplots.AxesSubplot at 0x23bdd5f6a20>
```



```
[95]: df_scaled = df1.apply(zscore)
```

```
[96]: df_scaled.head()
```

```
[96]:
```

	pelvic_incidence	pelvic_tilt numeric	lumbar_lordosis_angle	sacral_slope	\
0	0.147086	0.501369	-0.665177	-0.184950	
1	-1.245864	-0.748769	-1.453001	-1.041521	
2	0.484370	0.467932	-0.099262	0.273083	
3	0.511390	0.711562	-0.411339	0.126128	
4	-0.626648	-0.789693	-1.274745	-0.215876	

	pelvic_radius	degree_spondylolisthesis	class
0	-1.447647	-0.708059	0.690066
1	-0.264385	-0.579556	0.690066
2	-0.897686	-0.795421	0.690066
3	-1.207303	-0.402288	0.690066
4	-0.733455	-0.490106	0.690066

K-MEANS CLUSTERING ALGORITHM:

Finding the best K value:

```
[97]: model = KMeans(n_clusters = 3)
model
```

```
[97]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=3, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=None, tol=0.0001, verbose=0)
```

```
[184]: cluster_range = range( 1, 15 )
cluster_errors = []
for num_clusters in cluster_range:
    clusters = KMeans( num_clusters, n_init = 10 )
    clusters.fit(df_scaled)
    labels = clusters.labels_
    centroids = clusters.cluster_centers_
    cluster_errors.append( clusters.inertia_ )
clusters_df = pd.DataFrame( { "num_clusters":cluster_range, "cluster_errors":
    ↪cluster_errors } )
clusters_df[0:15]
```


```
[184]:
```

	num_clusters	cluster_errors
0	1	2170.000000
1	2	1390.613323
2	3	1144.373088
3	4	948.686227
4	5	819.948533
5	6	731.346852
6	7	661.204500
7	8	614.232600
8	9	573.813822
9	10	542.402881
10	11	514.985561
11	12	487.326659
12	13	472.131678
13	14	450.485634

```
[185]: # Elbow plot

plt.figure(figsize=(12,6))
plt.plot( clusters_df.num_clusters, clusters_df.cluster_errors, marker = "o" )
```

```
[185]: [<matplotlib.lines.Line2D at 0x23bd9ed87b8>]
```



```
[186]: df1.columns
```

```
[186]: Index(['pelvic_incidence', 'pelvic_tilt numeric', 'lumbar_lordosis_angle',
          'sacral_slope', 'pelvic_radius', 'degree_spondylolisthesis', 'class'],
          dtype='object')
```

here we are taking the value of k as 2

```
[187]: pi = sns.FacetGrid(df, col='class')
pi.map(sns.boxplot, 'pelvic_incidence', color='yellow', order=['0', '1'])

ptn= sns.FacetGrid(df, col='class')
ptn.map(sns.boxplot, 'pelvic_tilt numeric', color='orange', order=['0', '1'])


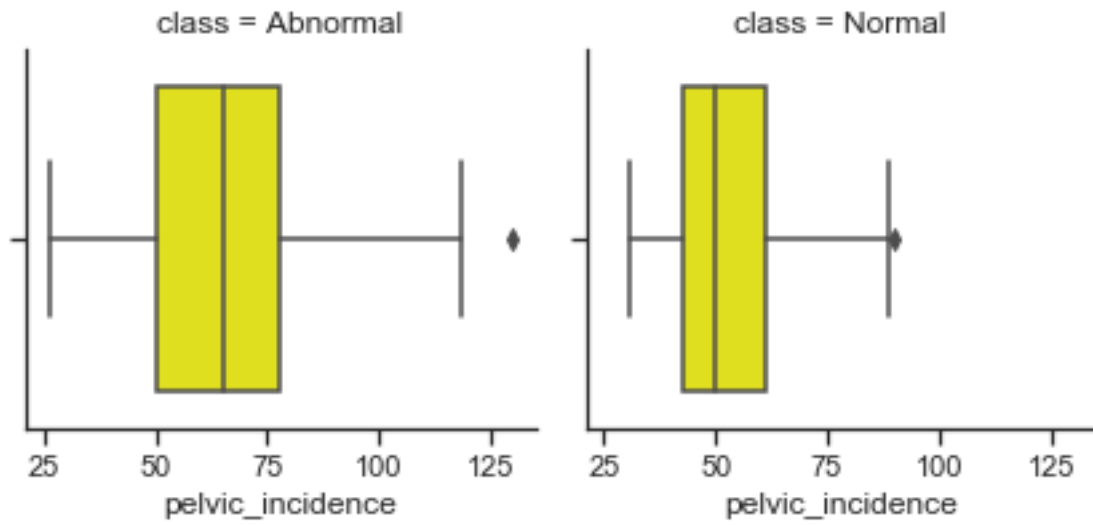
lla = sns.FacetGrid(df, col='class')
lla.map(sns.boxplot, 'lumbar_lordosis_angle', color='red', order=['0', '1'])

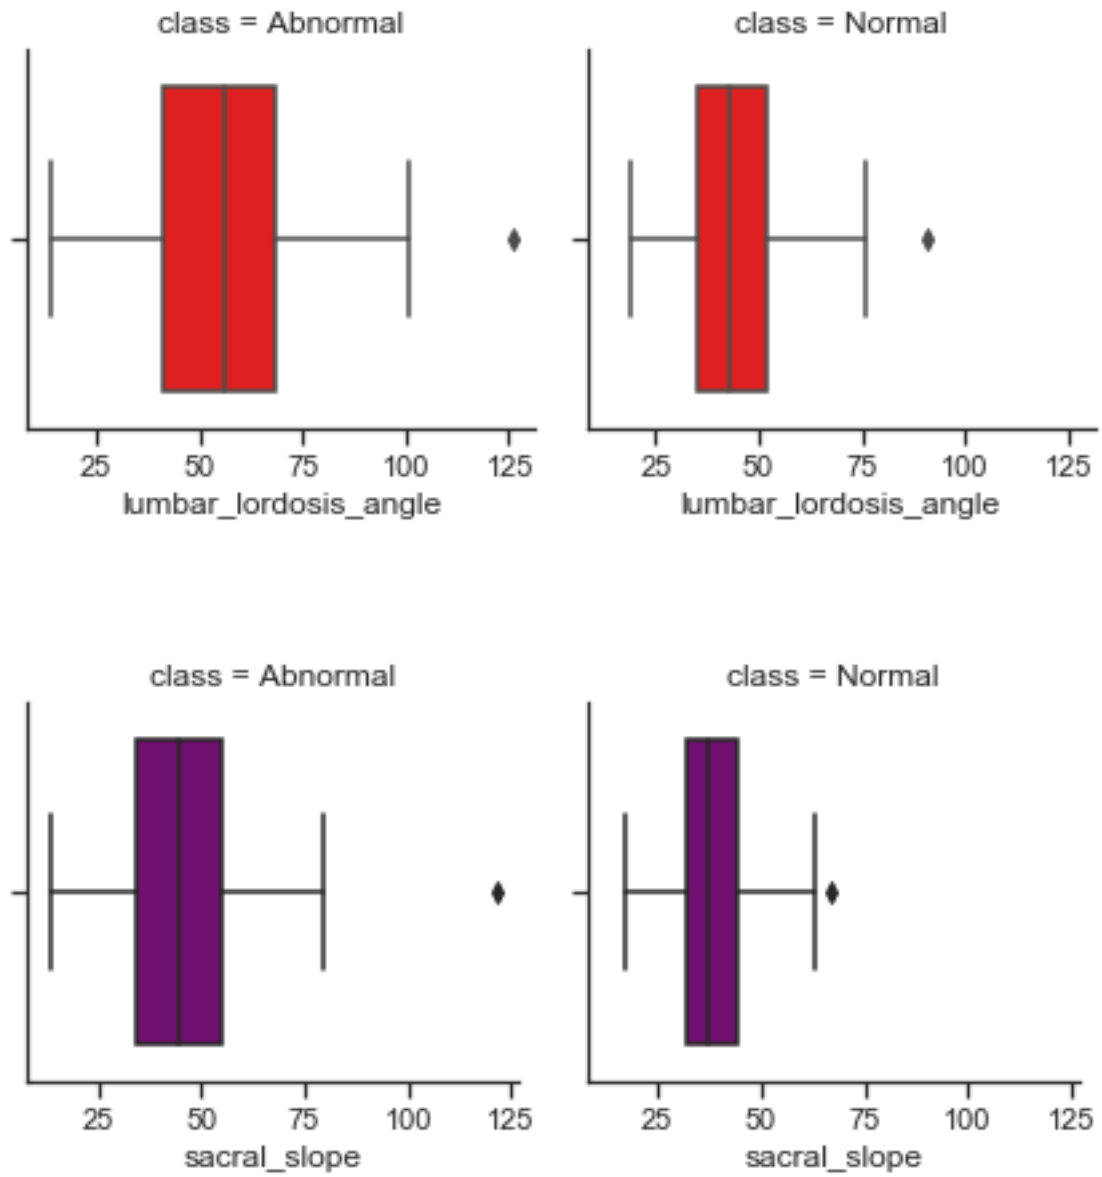
ss = sns.FacetGrid(df, col='class')
ss.map(sns.boxplot, 'sacral_slope', color='purple', order=['0', '1'])

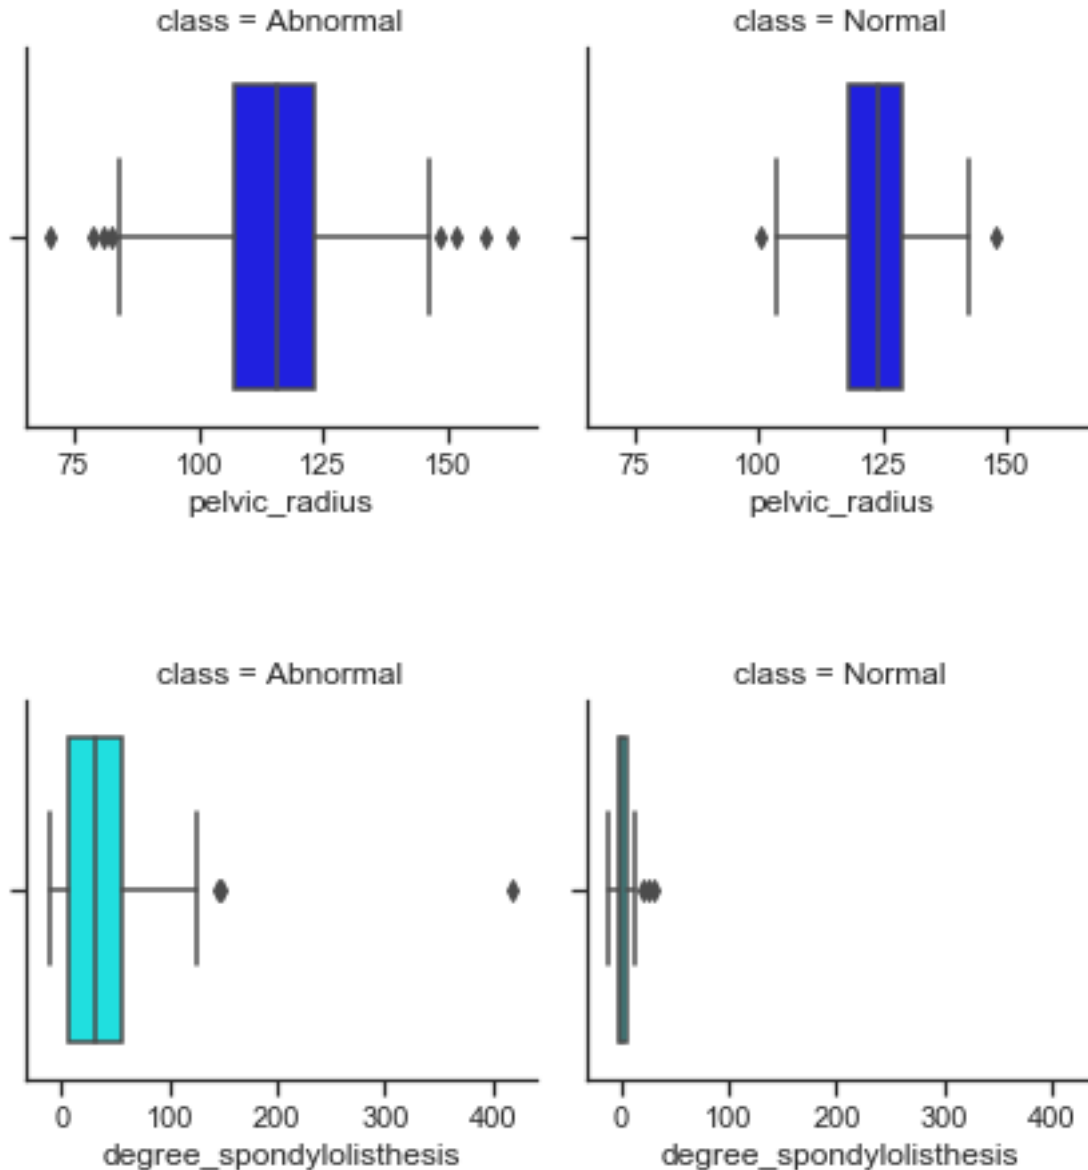
pr = sns.FacetGrid(df, col='class')
pr.map(sns.boxplot, 'pelvic_radius', color='blue', order=['0', '1'])

ds = sns.FacetGrid(df, col='class')
ds.map(sns.boxplot, 'degree_spondylolisthesis', color='cyan', order=['0', '1'])
```

```
[187]: <seaborn.axisgrid.FacetGrid at 0x23bd85a6128>
```





```
[188]: # Now we know our best k value is 2, I am creating a new kmeans model:
kmeans2 = KMeans(n_clusters=2)


# Training the model:
clusters = kmeans2.fit_predict(df1)

# Adding a label feature with the predicted class values:
df_k = df1.copy(deep=True)
df_k['label'] = clusters
```

```
[190]: fig, (ax1, ax2) = plt.subplots(1,2)

ax1 = plt.subplot(1,2,1)
plt.title('Original Classes')
sns.scatterplot(x='pelvic_incidence', y='pelvic_tilt numeric', hue='class',
               style='class', data=df1, ax=ax1)

ax2 = plt.subplot(1,2,2)
plt.title('Predicted Classes')
sns.scatterplot(x='pelvic_incidence', y='pelvic_tilt numeric', hue='label',
               style='label', data=df_k, ax=ax2)
plt.show()
```



```
[191]: print('Original Data Classes:')
print(df1['class'].value_counts())
print('-' * 30)
print('Predicted Data Classes:')
print(df_k.label.value_counts())
```

```
Original Data Classes:
1    210
0    100
Name: class, dtype: int64
-----
```

Predicted Data Classes:

1 200


0 110

Name: label, dtype: int64

HIERARCHICAL CLUSTERING ALGORITHM:

Creating the Dendrogram:

```
[192]: from scipy.cluster.hierarchy import linkage, dendrogram
plt.figure(figsize=[10,10])
merg = linkage(df_scaled, method='ward')
dendrogram(merg, leaf_rotation=90)
plt.title('Dendrogram')
plt.xlabel('Data Points')
plt.ylabel('Euclidean Distances')
plt.show()
```



From the dendrogram we can read there are 2 classes in our data set.

Hierarchical Clustering Algorithm:

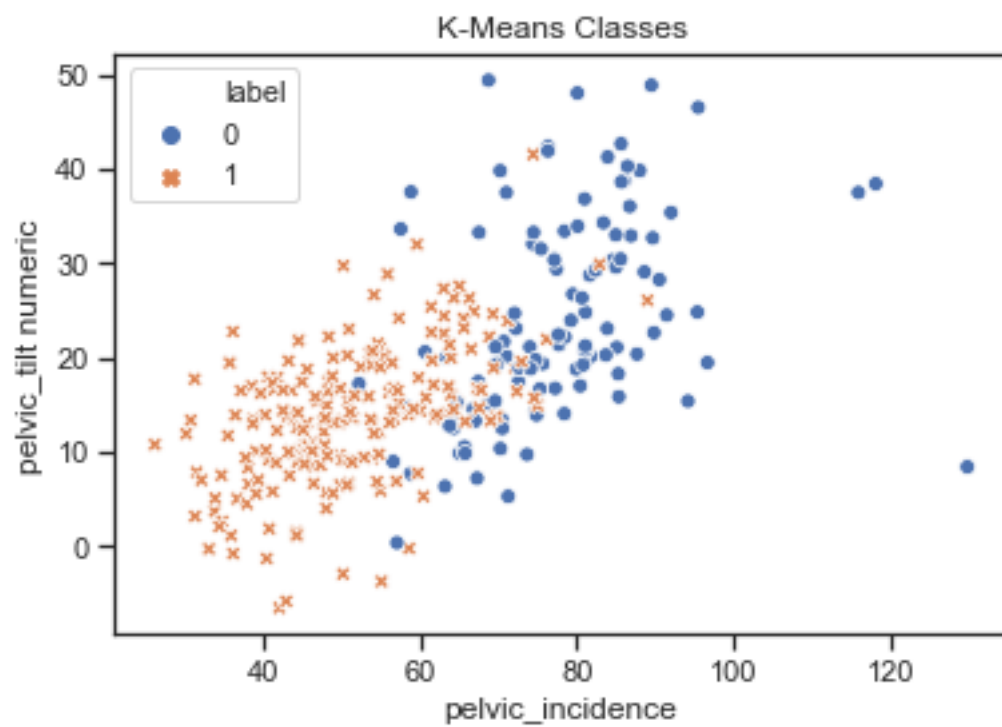

```
[193]: from sklearn.cluster import AgglomerativeClustering

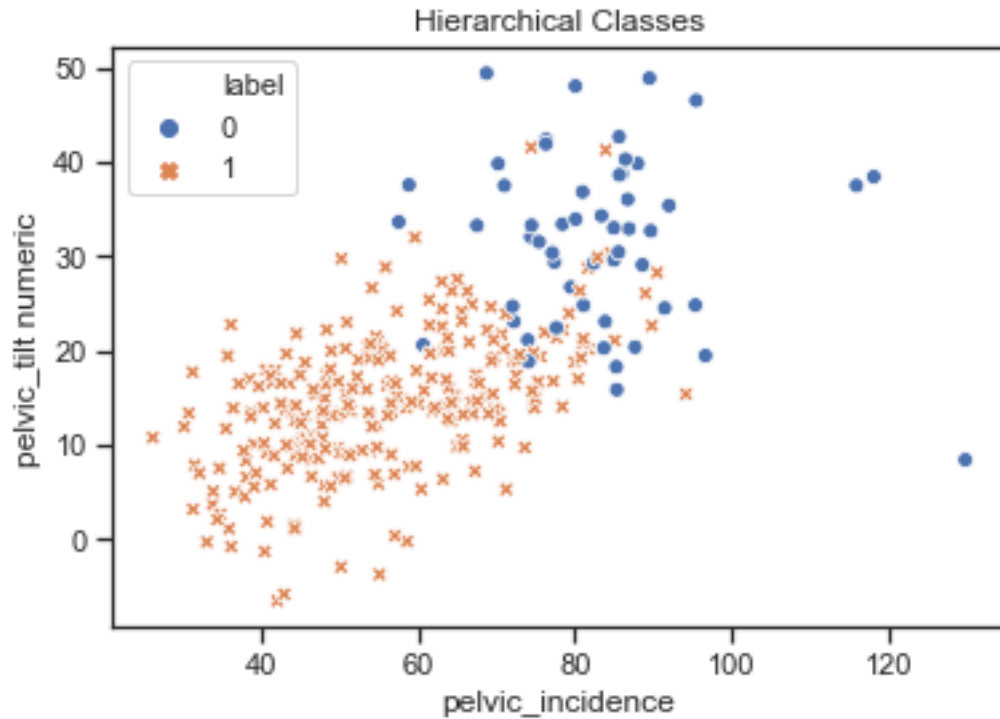
hie_clus = AgglomerativeClustering(n_clusters=2, affinity='euclidean',
    ↪linkage='ward')
cluster2 = hie_clus.fit_predict(df1)

df_h = df.copy(deep=True)
df_h['label'] = cluster2
```

Comparing Original, K-Means and Hierarchical Clustered Classes:

```
[194]: plt.title('Original Classes')
sns.scatterplot(x='pelvic_incidence', y='pelvic_tilt numeric', hue='class',
    ↪style='class', data=df1)
plt.show()
plt.title('K-Means Classes')
sns.scatterplot(x='pelvic_incidence', y='pelvic_tilt numeric', hue='label',
    ↪style='label', data=df_k)
plt.show()
plt.title('Hierarchical Classes')
sns.scatterplot(x='pelvic_incidence', y='pelvic_tilt numeric', hue='label',
    ↪style='label', data=df_h)
plt.show()
```





```
[195]: print('Original Data Classes:')
print(df1['class'].value_counts())
print('-' * 30)
print('K-Means Predicted Data Classes:')
print(df_k.label.value_counts())
print('-' * 30)
print('Hierarchical Predicted Data Classes:')
print(df_h.label.value_counts())
```

Original Data Classes:

1 210

0 100

Name: class, dtype: int64

K-Means Predicted Data Classes:

1 200

0 110

Name: label, dtype: int64

Hierarchical Predicted Data Classes:

1 256

0 54

Name: label, dtype: int64

We can see our models' differences from the comparison of our algorithms' class counts

2 Build An Classification model :-

3 K-Means

```
[196]: df_k.sample(5)
```

```
[196]:      pelvic_incidence  pelvic_tilt_numeric  lumbar_lordosis_angle \
133          81.754419          20.123466          70.560440
64           76.147212          21.936186          82.961502
200          63.364339          20.024621          67.498705
30           50.819268          15.402213          42.528939
227          61.540599          19.676957          52.892229

      sacral_slope  pelvic_radius  degree_spondylolisthesis  class  label
133      61.630954      119.425086          55.506889         1         0
64       54.211027      123.932010          10.431972         1         1
200      43.339718      130.999258          37.556706         1         0
30       35.417055      112.192804          10.869566         1         1
227      41.863642      118.686268           4.815031         0         1
```

Data Preparation(Splitting the Dependent/Target Variable and the Independent Variables)

```
[198]: x= df_k.drop('label',axis=1)
      y= df_k['label']
```

```
[199]: test_size = 0.30 # taking 70:30 training and test set
      seed = 7 # Random number seeding for repeatability of the code
      x_train, x_validate, y_train, y_validate = train_test_split(x, y,
      ↪test_size=test_size, random_state=seed)
```

```
[200]: from sklearn.preprocessing import StandardScaler
      independent_scalar = StandardScaler()
      x_train = independent_scalar.fit_transform(x_train) #fit and transform
      x_validate = independent_scalar.transform(x_validate) # only transform
```

Decision Tree Classifier

```
[201]: from sklearn.tree import DecisionTreeClassifier
      #DecisionTreeClassifier is the corresponding Classifier
      Dtree = DecisionTreeClassifier(max_depth=3)
      Dtree.fit(x_train, y_train)
```

```
[201]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
      max_depth=3, max_features=None, max_leaf_nodes=None,
      min_impurity_decrease=0.0, min_impurity_split=None,
```

```
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort='deprecated',
random_state=None, splitter='best')
```

```
[202]: predictValues_train = Dtree.predict(x_train)
        #print(predictValues_train)
        accuracy_train=accuracy_score(y_train, predictValues_train)

        predictValues_validate = Dtree.predict(x_validate)
        #print(predictValues_validate)
        accuracy_validate=accuracy_score(y_validate, predictValues_validate)

        print("Train Accuracy  :: ",accuracy_train)
        print("Validation Accuracy  :: ",accuracy_validate)
```

```
Train Accuracy  ::  0.9953917050691244
Validation Accuracy  ::  0.946236559139785
```

```
[203]: print('Classification Report')
        print(classification_report(y_validate, predictValues_validate))
```

```
Classification Report
```

	precision	recall	f1-score	support
0	0.97	0.89	0.93	38
1	0.93	0.98	0.96	55
accuracy			0.95	93
macro avg	0.95	0.94	0.94	93
weighted avg	0.95	0.95	0.95	93

Random Forest

```
[204]: RFclassifier = RandomForestClassifier(n_estimators = 100, random_state = 0,
        min_samples_split=5,criterion='gini',max_depth=5)
        RFclassifier.fit(x_train, y_train)
```

```
[204]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
        criterion='gini', max_depth=5, max_features='auto',
        max_leaf_nodes=None, max_samples=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=5,
        min_weight_fraction_leaf=0.0, n_estimators=100,
        n_jobs=None, oob_score=False, random_state=0, verbose=0,
        warm_start=False)
```

```
[205]: predictValues_validate = RFclassifier.predict(x_validate)
        #print(predictValues_validate)
        accuracy_validate=accuracy_score(y_validate, predictValues_validate)

        predictValues_train = RFclassifier.predict(x_train)
        #print(predictValues_train)
        accuracy_train=accuracy_score(y_train, predictValues_train)

        print("Train Accuracy  :: ",accuracy_train)
        print("Validation Accuracy  :: ",accuracy_validate)
```

```
Train Accuracy  ::  1.0
Validation Accuracy  ::  0.978494623655914
```

```
[206]: RFclassifier = RandomForestClassifier(n_estimators = 11, random_state = 0,
        min_samples_split=5,criterion='gini',max_depth=5)
        RFclassifier.fit(x_train, y_train)
```

```
[206]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
        criterion='gini', max_depth=5, max_features='auto',
        max_leaf_nodes=None, max_samples=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=5,
        min_weight_fraction_leaf=0.0, n_estimators=11,
        n_jobs=None, oob_score=False, random_state=0, verbose=0,
        warm_start=False)
```

```
[207]: predictValues_validate = RFclassifier.predict(x_validate)
        #print(predictValues_validate)
        accuracy_validate=accuracy_score(y_validate, predictValues_validate)

        predictValues_train = RFclassifier.predict(x_train)
        #print(predictValues_train)
        accuracy_train=accuracy_score(y_train, predictValues_train)

        print("Train Accuracy  :: ",accuracy_train)
        print("Validation Accuracy  :: ",accuracy_validate)
```

```
Train Accuracy  ::  0.9953917050691244
Validation Accuracy  ::  0.967741935483871
```

```
[208]: print('Classification Report')
print(classification_report(y_validate, predictValues_validate))
```

```
Classification Report
              precision    recall  f1-score   support

     0           1.00        0.92        0.96         38
     1           0.95        1.00        0.97         55

 accuracy              0.97
 macro avg              0.97
weighted avg              0.97
```

KNN

```
[121]: from sklearn.neighbors import KNeighborsClassifier
from scipy.stats import zscore
```

```
[143]: df_k
```

```
[143]:      pelvic_incidence  pelvic_tilt  numeric  lumbar_lordosis_angle  \
0          63.027818          22.552586          39.609117
1          39.056951          10.060991          25.015378
2          68.832021          22.218482          50.092194
3          69.297008          24.652878          44.311238
4          49.712859           9.652075          28.317406
..          ...          ...          ...
305         47.903565          13.616688          36.000000
306         53.936748          20.721496          29.220534
307         61.446597          22.694968          46.170347
308         45.252792           8.693157          41.583126
309         33.841641           5.073991          36.641233

      sacral_slope  pelvic_radius  degree_spondylolisthesis  class
0          40.475232          98.672917          -0.254400         1
1          28.995960          114.405425           4.564259         1
2          46.613539          105.985135          -3.530317         1
3          44.644130          101.868495          11.211523         1
4          40.060784          108.168725           7.918501         1
..          ...          ...          ...          ...
305         34.286877          117.449062          -4.245395         0
306         33.215251          114.365845          -0.421010         0
307         38.751628          125.670725          -2.707880         0
308         36.559635          118.545842           0.214750         0
309         28.767649          123.945244          -0.199249         0
```

[310 rows x 7 columns]

```
[144]: x= df_k.drop('class',axis=1)
      y= df_k['class']
```

```
[134]: x_standardize = x.apply(zscore)
```

```
[135]: #KNN only takes array as input hence it is important to convert dataframe to
      →array
      x1 = np.array(x_standardize)
      y1 = np.array(y)
```

```
[136]: test_size = 0.30 # taking 70:30 training and test set
      seed = 7 # Random number seeding for repeatability of the code
      x_train, x_validate, y_train, y_validate = train_test_split(x1, y1,
      →test_size=test_size, random_state=seed)
```

```
[137]: KNN = KNeighborsClassifier(n_neighbors= 8 , weights = 'uniform',
      →metric='euclidean')
      KNN.fit(x_train, y_train)
```

```
[137]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
      metric_params=None, n_jobs=None, n_neighbors=8, p=2,
      weights='uniform')
```

```
[138]: predictValues_train = KNN.predict(x_train)
      print(predictValues_train)
      accuracy_train=accuracy_score(y_train, predictValues_train)
      print("Train Accuracy :: ",accuracy_train)
```

```
[1 1 1 1 1 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 1 1 0 1 0 1 0 0 1 0 1 1 0 1 0 1
 0 1 0 0 1 1 0 1 1 1 1 1 0 1 1 1 0 0 1 0 1 1 0 1 1 0 1 1 1 1 1 0 0 0 1 1 0
 1 1 0 1 1 1 1 1 0 1 0 0 1 0 0 0 1 0 1 1 1 1 1 0 0 0 0 1 0 1 1 1 1 1 1 1 1
 0 1 0 1 1 1 1 1 0 1 1 1 0 1 0 0 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 1 0 1 1 0 1
 1 1 1 1 0 0 1 1 1 1 1 1 1 0 1 0 1 1 0 0 1 0 1 1 1 1 1 1 1 0 0 0 1 1 0 1 1
 0 0 1 1 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 0 1 0 1 1 1 1 1 0 1 1 1 1]
```

Train Accuracy :: 0.8387096774193549

```
[139]: predictValues_validate = KNN.predict(x_validate)
      print(predictValues_validate)
      accuracy_validate=accuracy_score(y_validate, predictValues_validate)
      print("Validation Accuracy :: ",accuracy_validate)
```

```
[0 1 1 1 1 0 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1
 1 1 1 1 0 0 0 1 1 1 0 1 0 1 0 1 1 1 1 1 1 1 1 0 1 1 1 0 1 0 0 1 1 0 1 1
 0 1 1 0 0 1 1 1 0 1 0 0 1 0 0 1 1 1 1]
```

Validation Accuracy :: 0.8279569892473119

```
[132]: df_k.drop('label', axis=1, inplace=True)
```

Agglomerative clustering

```
[141]: df_h.sample(5)
```

```
[141]:      pelvic_incidence  pelvic_tilt  numeric  lumbar_lordosis_angle  \
152          63.404481          14.115327          48.136806
70           72.560702          17.385191          52.000000
291          51.079833          14.209935          35.951229
22           63.073611          24.413803          54.000000
59           48.109236          14.930725          35.564683

      sacral_slope  pelvic_radius  degree_spondylolisthesis  class  label
152    49.289153    111.916007          31.784495  Abnormal    1
70     55.175511    119.193724          32.108537  Abnormal    1
291    36.869898    115.803711           6.905090   Normal    1
22     38.659808    106.424329          15.779697  Abnormal    1
59     33.178512    124.056452           7.947905  Abnormal    1
```

```
[145]: x= df_k.drop('class',axis=1)
      y= df_k['class']
```

```
[146]: test_size = 0.30 # taking 70:30 training and test set
      seed = 7 # Random numbmer seeding for reapeatability of the code
      x_train, x_validate, y_train, y_validate = train_test_split(x, y,
      ↪test_size=test_size, random_state=seed)
```

```
[147]: from sklearn.preprocessing import StandardScaler
      independent_scalar = StandardScaler()
      x_train = independent_scalar.fit_transform (x_train) #fit and transform
      x_validate = independent_scalar.transform (x_validate) # only transform
```

Decision Tree Classifier

```
[148]: from sklearn.tree import DecisionTreeClassifier
      #DecisionTreeClassifier is the corresponding Classifier
      Dtree = DecisionTreeClassifier(max_depth=3)
      Dtree.fit (x_train, y_train)
```

```
[148]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
      max_depth=3, max_features=None, max_leaf_nodes=None,
      min_impurity_decrease=0.0, min_impurity_split=None,
      min_samples_leaf=1, min_samples_split=2,
      min_weight_fraction_leaf=0.0, presort='deprecated',
      random_state=None, splitter='best')
```

```
[149]: predictValues_train = Dtree.predict(x_train)
      #print(predictValues_train)
      accuracy_train=accuracy_score(y_train, predictValues_train)
```

```

predictValues_validate = Dtree.predict(x_validate)
#print(predictValues_validate)
accuracy_validate=accuracy_score(y_validate, predictValues_validate)

print("Train Accuracy  :: ",accuracy_train)
print("Validation Accuracy  :: ",accuracy_validate)

```

```

Train Accuracy  ::  0.8847926267281107
Validation Accuracy  ::  0.7741935483870968

```

```

[150]: print('Classification Report')
print(classification_report(y_validate, predictValues_validate))

```

```

Classification Report

```

	precision	recall	f1-score	support
0	0.81	0.50	0.62	34
1	0.76	0.93	0.84	59
accuracy			0.77	93
macro avg	0.79	0.72	0.73	93
weighted avg	0.78	0.77	0.76	93

Random Forest

```

[151]: RFclassifier = RandomForestClassifier(n_estimators = 100, random_state = 0,
min_samples_split=5,criterion='gini',max_depth=5)
RFclassifier.fit(x_train, y_train)

```

```

[151]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
criterion='gini', max_depth=5, max_features='auto',
max_leaf_nodes=None, max_samples=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=5,
min_weight_fraction_leaf=0.0, n_estimators=100,
n_jobs=None, oob_score=False, random_state=0, verbose=0,
warm_start=False)

```

```

[152]: predictValues_validate = RFclassifier.predict(x_validate)
#print(predictValues_validate)
accuracy_validate=accuracy_score(y_validate, predictValues_validate)

```

```

predictValues_train = RFclassifier.predict(x_train)
#print(predictValues_train)
accuracy_train=accuracy_score(y_train, predictValues_train)

print("Train Accuracy  :: ",accuracy_train)
print("Validation Accuracy  :: ",accuracy_validate)

```

Train Accuracy :: 0.967741935483871
Validation Accuracy :: 0.8709677419354839

```

[153]: RFclassifier = RandomForestClassifier(n_estimators = 11, random_state = 0,
min_samples_split=5,criterion='gini',max_depth=5)
RFclassifier.fit(x_train, y_train)

```

```

[153]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
criterion='gini', max_depth=5, max_features='auto',
max_leaf_nodes=None, max_samples=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=5,
min_weight_fraction_leaf=0.0, n_estimators=11,
n_jobs=None, oob_score=False, random_state=0, verbose=0,
warm_start=False)

```

```

[154]: predictValues_validate = RFclassifier.predict(x_validate)
#print(predictValues_validate)
accuracy_validate=accuracy_score(y_validate, predictValues_validate)

predictValues_train = RFclassifier.predict(x_train)
#print(predictValues_train)
accuracy_train=accuracy_score(y_train, predictValues_train)

print("Train Accuracy  :: ",accuracy_train)
print("Validation Accuracy  :: ",accuracy_validate)

```

Train Accuracy :: 0.9447004608294931
Validation Accuracy :: 0.9032258064516129

```

[155]: print('Classification Report')
print(classification_report(y_validate, predictValues_validate))

```

```

Classification Report
      precision    recall  f1-score   support

0             1.00      0.74      0.85         34

```


1	0.87	1.00	0.93	59
accuracy			0.90	93
macro avg	0.93	0.87	0.89	93
weighted avg	0.92	0.90	0.90	93

```
[156]: from sklearn.neighbors import KNeighborsClassifier
from scipy.stats import zscore
```

```
[158]: x= df_k.drop('class',axis=1)
y= df_k['class']
```

```
[159]: x_standardize = x.apply(zscore)
```

```
[160]: #KNN only takes array as input hence it is important to convert dataframe to
↳array
x1 = np.array(x_standardize)
y1 = np.array(y)
```

```
[161]: test_size = 0.30 # taking 70:30 training and test set
seed = 7 # Random number seeding for repeatability of the code
x_train, x_validate, y_train, y_validate = train_test_split(x1, y1,
↳test_size=test_size, random_state=seed)
```

```
[162]: predictValues_train = KNN.predict(x_train)
print(predictValues_train)
accuracy_train=accuracy_score(y_train, predictValues_train)
print("Train Accuracy :: ",accuracy_train)
```

```
[1 1 1 1 1 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 1 1 0 1 0 1 0 0 1 0 1 1 0 1 0 1
0 1 0 0 1 1 0 1 1 1 1 1 0 1 1 1 0 0 1 0 1 1 0 1 1 0 1 1 1 1 1 0 0 0 1 1 0
1 1 0 1 1 1 1 1 0 1 0 0 1 0 0 0 1 0 1 1 1 1 1 0 0 0 0 1 0 1 1 1 1 1 1 1 1
0 1 0 1 1 1 1 1 0 1 1 1 0 1 0 0 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 1 0 1 1 0 1
1 1 1 1 0 0 1 1 1 1 1 1 1 0 1 0 1 1 0 0 1 0 1 1 1 1 1 1 1 1 0 0 0 1 1 0 1 1
0 0 1 1 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 0 1 0 1 1 1 1 1 1 0 1 1 1 1]
```

Train Accuracy :: 0.8387096774193549

```
[163]: predictValues_validate = KNN.predict(x_validate)
print(predictValues_validate)
accuracy_validate=accuracy_score(y_validate, predictValues_validate)
print("Validation Accuracy :: ",accuracy_validate)
```

```
[0 1 1 1 1 0 1 1 0 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1
1 1 1 1 0 0 0 1 1 1 0 1 0 1 0 1 1 1 1 1 1 1 1 0 1 1 1 0 1 0 0 1 1 0 1 1
0 1 1 0 0 1 1 1 0 1 0 0 1 0 0 1 1 1 1]
```

Validation Accuracy :: 0.8279569892473119

[]: